

SPACEX

Space Exploration Technologies Corp.
Class A Common Stock

Filed Pursuant to Rule 424(b)(4)
Registration No. 333-296070

555,555,555 Shares

This is the initial public offering of shares of Class A common stock, par value \$0.001 per share, of Space Exploration Technologies Corp., a Texas corporation. We are offering 555,555,555 shares of our Class A common stock.

Currently, no public market exists for our Class A common stock. The initial public offering price is \$135.00 per share. We have been approved to list our Class A common stock on The Nasdaq Stock Market LLC (“Nasdaq”) and Nasdaq Texas, LLC (“Nasdaq Texas”) under the symbol “SPCX.”

Following the completion of this offering, we will have two classes of common stock issued and outstanding: Class A common stock and Class B common stock. Each share of Class A common stock will entitle its holder to one vote per share. Each share of Class B common stock will entitle its holder to 10 votes per share. Class A shareholders and Class B shareholders will vote together as a single class on all matters to be voted on by shareholders, except Class B shareholders will be entitled to elect a majority of our board of directors in addition to having certain other class votes as described under “Description of Capital Stock.”

Based on an offering size as set forth above and an initial public offering price of \$135.00 per share, Elon Musk, our founder, Chief Executive Officer, Chief Technical Officer and Chairman of our board, will hold approximately 82.4% of the voting power of our common stock (or approximately 82.3% if the underwriters exercise their option to purchase additional shares of Class A common stock in full) immediately after the completion of this offering through his ownership of shares of our Class A and Class B common stock, of which approximately 81.1% is attributable to his ownership of our Class B common stock. As a result, Mr. Musk will be able to control the outcome of matters requiring shareholder approval. This includes the election of (i) a majority of our board, through his ownership of Class B shares (as Class B Directors), for so long as he holds a majority of the voting power of the Class B common stock, and (ii) the remainder of our board, for so long as he holds a majority of the combined voting power of the Class A and Class B common stock. As a result, we will be a “controlled company” under the corporate governance rules of Nasdaq following the completion of this offering and, as a result, we intend to rely on exemptions from certain corporate governance requirements. Please refer to “Management—Controlled Company Exemption.”

Investing in our Class A common stock involves risks. Please refer to “Risk Factors” beginning on page 27 of this prospectus.

	Per Share	Total
Initial public offering price	\$ 135.00	\$ 74,999,999,925
Underwriting discounts and commissions ⁽¹⁾	\$ 0.90	\$ 500,000,000
Proceeds, before expenses, to Space Exploration Technologies Corp.	\$ 134.10	\$ 74,499,999,925

⁽¹⁾ The underwriters will not receive any discount or commission on any shares of our Class A common stock sold pursuant to the over-allotment option. Please refer to “Underwriting” for a description of all underwriting compensation payable in connection with this offering.

The underwriters may also exercise an option to purchase up to an additional 83,333,333 shares of our Class A common stock from us, at the initial public offering price for 30 days after the date of this prospectus.

At our request, the underwriters have reserved up to five percent of the shares of Class A common stock to be issued by the Company and offered by this prospectus for sale, at the initial public offering price, to certain employees and persons identified by our executive officers. Please refer to “Underwriting—Directed Share Program.” Neither the Securities and Exchange Commission (the “SEC”) nor any state securities commission has approved or disapproved of these securities or passed on the adequacy or accuracy of this prospectus. Any representation to the contrary is a criminal offense.

The shares of Class A common stock will be ready for delivery on or about June 15, 2026.

Joint Book-Running Managers

Goldman Sachs & Co. LLC

Morgan Stanley

BofA Securities

Citigroup

J.P. Morgan

Barclays

Deutsche Bank Securities

RBC Capital Markets

UBS Investment Bank

Wells Fargo Securities

Cantor

Needham & Company

Raymond James

Societe Generale

Stifel

William Blair

BTG Pactual

ING

Macquarie Capital

Mirae Asset Securities

Mizuho

Santander

Prospectus Dated June 11, 2026.





■ SPACE

ACTIVATED
2002

FIRST:

2008

Liquid fuel rocket to reach orbit by a private company

2012

Private company to successfully dock a spacecraft with the International

■ CONNECTIVITY

ACTIVATED
2020

FIRST:

2019

To begin deploying a large-scale low-Earth orbit broadband satellite constellation

2022

Low-latency low-Earth orbit network available globally

2022

To manufacture consumer-grade phased-array user terminals at scale

2025

To deploy a large-scale low-Earth orbit satellite-to-mobile constellation

HIGHLIGHTS

Starlink
Broadband &
Mobile
Satellites

9,600+

% of Active
Maneuverable
Satellites

~75%

Starlink
Subscribers

~10.3M

Countries

164

Highlights as of March 31, 2026. Please refer to "Management's Discussion and Analysis of Financial Condition and Results of Operation-Key Business Metrics" for information about the key business metrics our management uses to evaluate our performance.

■ AI

ACTIVATED
2023

FIRST:

2026

Gigawatt-scale AI training cluster

2026

Gigawatt-scale Megapack Battery
Installation

HIGHLIGHTS

Nameplate
Compute Draw

1GW+


Major Model
Versions


4

















[Table of Contents](#)

TABLE OF CONTENTS

GLOSSARY OF TERMS	iv
PROSPECTUS SUMMARY	1
RISK FACTORS	27
CAUTIONARY STATEMENT REGARDING FORWARD-LOOKING STATEMENTS	65
USE OF PROCEEDS	67
DIVIDEND POLICY	68
CAPITALIZATION	69
DILUTION	71
MANAGEMENT'S DISCUSSION AND ANALYSIS OF FINANCIAL CONDITION AND RESULTS OF OPERATIONS	74
BUSINESS	130
MANAGEMENT	226
EXECUTIVE COMPENSATION	233
CERTAIN RELATIONSHIPS AND RELATED PERSON TRANSACTIONS	243
SECURITY OWNERSHIP OF CERTAIN BENEFICIAL OWNERS AND MANAGEMENT	247
DESCRIPTION OF CAPITAL STOCK	250
SHARES ELIGIBLE FOR FUTURE SALE	258
MATERIAL U.S. FEDERAL INCOME TAX CONSIDERATIONS FOR NON-U.S. HOLDERS OF CLASS A COMMON STOCK	262
UNDERWRITING	266
LEGAL MATTERS	283

Through and including July 6, 2026 (the 25th day after the date of this prospectus), all dealers effecting transactions in our Class A common stock, whether or not participating in this offering, may be required to deliver a prospectus. This delivery requirement is in addition to a dealer's obligation to deliver a prospectus when acting as an underwriter and with respect to an unsold allotment or subscription.

Neither we nor the underwriters have authorized anyone to provide you with information other than that contained in this prospectus or in any free writing prospectus authorized by us. We and the underwriters take no responsibility for, and can provide no assurance as to the reliability of, any other information that others may give you. We and the underwriters are not making an offer to sell, or seeking offers to buy, our Class A common stock in any jurisdiction where an offer or sale is not permitted. The information contained in this prospectus or any free writing prospectus is accurate only as of its date, regardless of its time of delivery or of any sale of shares of our Class A common stock. Our business, financial condition, results of operations and future prospects may have changed since that date.

For investors outside of the United States: Neither we nor the underwriters have done anything that would permit this offering, or possession or distribution of this prospectus, in any jurisdiction where action for that purpose is required, other than the United States. Persons outside of the United States who come into possession of this prospectus must inform themselves about, and observe any restrictions relating to, the offering of the shares of our Class A common stock and the distribution of this prospectus outside of the United States.

This prospectus contains forward-looking statements that are subject to a number of risks and uncertainties, many of which are beyond our control. Please refer to "Risk Factors" and "Cautionary Statement Regarding Forward-Looking Statements."

[Table of Contents](#)

General Information

Except as otherwise indicated or required by the context, all references to "SpaceX," the "Company," "we," "our" and "us" or similar terms refer to Space Exploration Technologies Corp. and its consolidated subsidiaries. For the definitions of certain terms and abbreviations used in this prospectus, please refer to "Glossary of Terms" beginning on page [iv](#) of this prospectus.

References to (i) our "bylaws" are to the form of amended and restated bylaws of the Company (as amended and restated from time to time) to be effective upon the completion of this offering, (ii) our "charter" are to the form of restated certificate of formation of the Company to be effective upon the completion of this offering and (iii) "our board" or "the board" are to the board of directors of the Company.

Basis of Presentation

The consolidated financial statements of SpaceX have been retrospectively recast for all periods presented to include (i) the historical results of X.AI Holdings Corp., which was acquired by SpaceX, effective February 2, 2026 (the "xAI Merger"), and X Holdings Corp. ("X Holdings"), which was acquired by xAI, effective March 28, 2025 (the "X Merger"), because these transactions were between entities under common control, and (ii) a five-for-one stock split of the Company's Class A, Class B, and Class C Common Stock, effective May 4, 2026 (the "2026 Stock Split"). Unless otherwise stated or the context otherwise requires, all share and per share information included in this prospectus have been retroactively adjusted to reflect the 2026 Stock Split. Refer to Note 1, Nature of Business, to the audited consolidated financial statements included elsewhere in this prospectus.

Industry and Market Data

Certain market and industry data and forecasts used in this prospectus have been obtained from, are based on, or use data from, the following reports and sources, among others: (i) *Breaking Barriers to Data Center Growth*, dated January 20, 2025, by Boston Consulting Group; (ii) *Looming Spectrum Shortfall Could Cost America's GDP \$1.4T, Jeopardize Continued Function of U.S. Networks, New Report Finds*, dated March 27, 2025, by the Cellular Telecommunications and Internet Association; (iii) *Top 50 Countries by Number of Business Aircraft Registered*, dated January 27, 2026, by Corporate Jet Investor; (iv) *Digital Economy Trends 2026*, dated December 2025, by the Digital Cooperation Organization; (v) *Global Fixed Broadband Market Outlook*, Ericsson Mobility Report, dated November 1, 2025, by Ericsson; (vi) *Households by Number of Households and by Country*, Euromonitor International Passport 2026 Edition, dated November 5, 2025, by Euromonitor International; (vii) *Satellite Solutions for Universal Service*, dated March 2025, by the Global Satellite Operators Association; (viii) *Broadband Services Market Analysis Segment Forecast to 2027*, dated April 2025, by Grand View Research; (ix) *Consumer Market Model H2 2025 – Worldwide Household Internet Penetration*, dated March 2026, by International Data Corporation; (x) *World Energy Outlook Special Report: Energy and AI*, dated April 2025, by the International Energy Agency; (xi) *The 175 GW Crisis: America's Power Grid Cannot Keep Up with AI Data Centers*, dated January 21, 2026, by Introl; (xii) *As Wireless Network Quality Competition Increases, Customers Benefit*, dated July 17, 2025, by J.D. Power; (xiii) *Satellite Statistics: Satellite and Debris Population*, dated April 2026, by Jonathan McDowell; (xiv) *2026 Global Data Center Outlook: Navigating AI Demand, Power Constraints and Global Opportunities*, dated January 5, 2026, by JLL; (xv) *Global Ship Tracking Intelligence*, at marinetraffic.com, as updated from time to time and last accessed April 13, 2026, by Marine Traffic Dashboard; (xvi) *The Cost of Compute: A \$7 Trillion Race to Scale Data Centers*, dated April 28, 2025, by McKinsey & Company; (xvii) *What is Multimodal AI?*, dated June 10, 2025, by McKinsey & Company; (xviii) *NASA: Enabling America on the Space Frontier*, dated December 2024, by the National Aeronautics and Space Administration ("NASA"); (xix) *Space Act Agreement*, dated April 2015, by NASA; (xx) *The Recent Large Reduction in Space Launch Cost*, dated July 8, 2018, by NASA; (xxi) *12th Edition Space Economy Report*, dated January 29, 2026, by Novaspaces; (xxii) *Global Fleet and MRO Market Forecast 2025–2035*, dated February 2025, by Oliver Wyman; (xxiii) *Broadband Op Subs by Technology – Forecasts Summary*, dated March 31, 2026, by Omdia; (xxiv) *Mobile Forecasts Summary – February 2026*, dated February 18, 2026, by Omdia; (xxv) *Data Center Rules and Regulations*, dated September 8, 2025, by QTS; (xxvi) *AI's Power Requirements Under Exponential Growth*, dated January 28, 2025, by RAND Corporation; (xxvii) *Data Center Grid-Power Demand to Rise 22% in 2025, Nearly Triple by 2030*, dated October 14, 2025, by S&P Global Market Intelligence; (xxviii) *NVIDIA GTC 2025 – Built for Reasoning*, Vera Rubin, Kyber, CPO, Dynamo

[Table of Contents](#)

Inference, Jensen Math, Feynman, dated March 18, 2025, by SemiAnalysis; (xxix) *NVIDIA Blackwell Ultra Datasheet*, dated February 16, 2026, by SemiAnalysis; (xxx) *H100 Rental Price Over Time (2023–2025): A Complete Market Analysis*, dated December 21, 2025, by Silicon Data; (xxxi) *Data Centers – Understanding the Power Consumption of Data Centers*, at socomec.us, as updated from time to time and last accessed April 13, 2026, by Socomec; (xxxii) *The Space Report 2025 Q2 Highlights Record \$613 Billion Global Space Economy for 2024*, dated July 22, 2025, by the Space Foundation; (xxxiii) *Median Country Speeds Updated February 2026*, dated February 2026, by the Speedtest Global Index; (xxxiv) *Data Center (Russian Market) Commercial Data Centers*, dated January 28, 2026, by TAdviser; (xxxv) *Merchant Fleet by Flag of Registration and by Type of Ship*, dated June 10, 2025, by the United Nations Conference on Trade and Development; (xxxvi) *U.S. Electricity Generation in 2025 Hit a Record, Again*, dated March 5, 2026, by the U.S. Energy Information Administration; (xxxvii) *GAO-25-107555, In-Space Servicing, Assembly, and Manufacturing: Benefits, Challenges, and Policy Options*, dated July 2025, by the U.S. Government Accountability Office; (xxxviii) *GDP (current US\$)*, at data.worldbank.org, as updated from time to time and last accessed April 13, 2026, by the World Bank; (xxxix) *Rural population (% of total population)*, at data.worldbank.org, as updated from time to time and last accessed May 2, 2026, by the World Bank; (xl) *How Data Centres in Space Sustainably Enable the AI Revolution*, dated January 16, 2026, by Philip Johnston Co-Founder and Chief Executive Officer, Starcloud, published by the World Economic Forum; and (xli) *Most Americans Use AI but Still Don't Trust It*, dated December 9, 2025, by YouGov. We did not commission the preparation of any of these reports or sources.

Some market data and statistical information contained in this prospectus are also based on management's estimates and calculations, which are derived from our review and interpretation of publicly available industry publications, our internal research and our knowledge of the markets in which we currently, and will in the future, operate, as well as the sources referred to above. This information involves a number of assumptions and limitations, and you are cautioned not to give undue weight to such information. The estimates and assumptions used in determining our total addressable markets are further detailed in the section titled "Business—Our Market Opportunity," and you are urged to read the risk factor titled "The estimates of future market opportunity and forecasts of market growth, and our ability to capture such markets, included in this prospectus may prove to be inaccurate." Forecasts and other forward-looking information obtained from the sources named above are subject to the same qualifications and uncertainties as the other forward-looking statements in this prospectus.

Statements as to market position, market opportunity and market size are based on data currently available to us, as well as management's estimates, judgments, assessments, and assumptions. While we are not aware of any misstatements regarding market position, market opportunity, and market size information included in this prospectus, such information, which is derived in part from management's estimates and beliefs, is inherently uncertain and imprecise. Projections, assumptions and estimates of estimated market position and market opportunity and the future performance of the industries in which we operate are necessarily subject to a high degree of uncertainty and risk due to a variety of factors, including those described in "Risk Factors," "Cautionary Statement Regarding Forward-Looking Statements" and elsewhere in this prospectus. These and other factors could cause results to differ materially from those expressed in the estimates made by third parties and by us. Investors are cautioned not to place undue reliance on statements of expected future market size or opportunity.

Trademarks and Trade Names

We own or have rights to various trademarks, service marks and trade names that we use in connection with the operation of our business. This prospectus may also contain trademarks, service marks and trade names of third parties, which are the property of their respective owners. Our use or display of third parties' trademarks, service marks, trade names or products in this prospectus is not intended to, and does not imply, a relationship with us or an endorsement or sponsorship by or of us. Solely for convenience, the trademarks, service marks and trade names referred to in this prospectus may appear without the ®, ™ or SM symbols, but such references are not intended to indicate, in any way, that we will not assert, to the fullest extent under applicable law, our rights or the right of the applicable licensor to these trademarks, service marks and trade names.

[Table of Contents](#)

GLOSSARY OF TERMS

The terms and abbreviations defined in this section are used throughout this prospectus:

- "AI" or "artificial intelligence" refers to advanced computational technologies and systems enabling machines to learn, comprehend reality, solve complex problems, exhibit creativity, make critical decisions, and function with growing autonomy.
- "AI compute" or "compute" refers to the computing infrastructure required to train and operate artificial intelligence models, including, without limitation, specialized processors, networking, storage, and power systems deployed in data centers or other computing environments.
- "AI compute satellite" refers to a satellite equipped with onboard artificial intelligence processing capabilities designed to perform data analysis, inference, or other machine learning, automated decision-making and artificial intelligence algorithms, models and technologies workloads in orbit.
- "AI ecosystem" refers to a complex, multi-layered network of technologies, products, systems, and

infrastructure that develop, leverage, and deploy intelligent systems.

- “AI segment” refers to our AI business, which we acquired in connection with our acquisition of xAI in February 2026, and includes our AI compute, Grok, and X.
- “AI training cluster” refers to an integrated system that provides computational power required for training and running advanced AI models.
- “The Algorithm” refers to our five-step iterative process that we use to rapidly innovate and optimize, emphasizing making the requirements less dumb, deleting unnecessary processes or parts, optimizing the necessary processes or parts, accelerating cycle timesteps, and automating only proven processes after the first four steps are completed.
- “Application Programming Interface” or “API” refers to a defined set of rules and protocols that allows different software systems to communicate with and interact with each other programmatically.
- “ARPU” refers to service revenue generated from Starlink Subscribers during a period divided by (i) the average number of Starlink Subscribers during the period and by (ii) the number of months in the period.
- “Artemis program” refers to a NASA program aimed at landing humans on the Moon by the late 2020s.
- “booster” refers to the first-stage rocket that provides the primary thrust during launch.
- “booster catch” refers to a recovery method in which a returning first-stage rocket booster is captured mid-air by mechanical arms on the launch tower rather than on legs at a landing zone or at sea.
- “booster launch” refers to a rocket launch in which a booster stage provides the primary thrust during liftoff and the initial phase of ascent before separating from the vehicle.
- “bps” refers to bits per second.
- “COLOSSUS” refers to our flagship data center, located on Paul R. Lowry Road in Memphis, Tennessee.
- “COLOSSUS II” refers to our data centers in Memphis, Tennessee and in Southaven, Mississippi. These data centers are part of our coherent gigawatt-scale AI training cluster.
- “Connectivity segment” refers to our Connectivity segment, which includes Starlink and associated offerings.
- “Credit Agreements” refers to our SpaceX Credit Facility and SpaceX Bridge Loan.
- “crewmember” refers to a person who has traveled on our spacecraft, measuring by each mission.

[Table of Contents](#)

- “daily posts” on X and Grok refers to the aggregate volume of original posts, replies, reposts, quotes and media shared daily by users on the X platform, and the real-time interactions, analysis and generative capabilities provided to a user by Grok. This may include posts generated by AI or accounts managed by AI.
- “downlink capacity” refers to the maximum rate at which data can be transmitted from a satellite to users over a network or communication link in a given period of time.
- “Draco thrusters” refers to thrusters used in Dragon spacecraft for precise orbital maneuvering and adjustments.
- “Dragon” refers to our Dragon spacecraft.
- “Falcon 1” refers to our two-stage, liquid-fueled small-lift launch vehicle that operated from 2006 to 2009.
- “Falcon 9” refers to our orbital-class rocket with reusable boosters, first launched in 2010, which has a payload capacity to LEO of approximately 23 metric tons.
- “Falcon Heavy” refers to our partially reusable super heavy-lift launch vehicle, first launched in 2018, which has a payload capacity to LEO of approximately 64 metric tons.
- “flight-proven booster launches” refers to a mission utilizing a booster that has previously completed at least one successful launch and recovery.
- “frontier model” refers to a leading-edge, sophisticated large language model, such as Grok, designed for rigorous reasoning and real-time information synthesis.
- “Gbps” refers to gigabits per second.
- “geostationary orbit” refers to a high Earth orbit that allows satellites to match Earth’s rotation, appearing stationary from the ground, often used for communication satellites.
- “geosynchronous transfer orbit” refers to an elliptical orbit used to transfer a spacecraft from a lower orbit to a geostationary orbit.
- “gigawatt” refers to one billion watts.
- “gigawatt-scale” refers to infrastructure, systems, or facilities that are designed to generate, transmit, or consume approximately one gigawatt or more of electrical power capacity.
- “GPU” refers to a graphics processing unit.
- “Grok” refers to our family of frontier models, which represents a core pillar of our mission to advance humanity’s understanding of the universe through the development of truth-seeking artificial intelligence.
- “Grok API” refers to our application programming interface that enables developers to access and integrate Grok models into external software applications and workflows.
- “Grok Business” refers to our subscription-based offering that provides organizations with access to Grok

models and related tools for use in internal business applications and workflows, designed for deployment by small-to-medium teams.

- “Grok Enterprise” refers to our subscription-based offering that provides organizations with access to Grok models and related tools for use in internal business applications and workflows, designed for deployment by enterprise organizations.
- “Grok Voice” refers to the Grok real-time speech engine.
- “high-density compute” refers to compute infrastructure designed to deliver a large amount of processing power within a limited physical footprint, typically characterized by high processor concentration and elevated power usage per unit of space.

[Table of Contents](#)

- “Imagine” refers to our image and video generation system.
- “inference” refers to the process by which a trained artificial intelligence model generates outputs (such as text, images, or predictions) from new input data.
- “International Docking System Standard” refers to a standard for autonomous docking capabilities used by spacecraft like Dragon.
- “IoT” refers to the network of physical objects embedded with sensors, software, and other technologies for the purpose of connecting and exchanging data with other devices and systems over the internet.
- “Kardashev Type II” refers to a civilization that harnesses the full energy output of its local star, like our Sun, to power unprecedented growth and sustain the civilization’s existence.
- “large language model” or “LLM” refers to a sophisticated artificial intelligence model designed for advanced reasoning and natural language processing.
- “large-scale LEO broadband satellite constellation” refers to a satellite constellation network of over 1,000 satellites.
- “latency” refers to the time delay between the transmission of data from a source and its receipt at a destination, typically measured in milliseconds.
- “launch payload mass” refers to the theoretical payload mass that a particular spacecraft is capable of delivering to a specified orbit under specific conditions, which is derived from advanced computer simulations and performance modeling that apply to particular mission scenarios and trajectory assumptions. Actual payload that can be delivered for a given mission may be different and will vary depending on numerous mission parameters and operational factors, including mission-specific trajectory requirements, atmospheric conditions, vehicle and payload configuration, risk profile, and applicable regulatory or range-safety limitations.
- “launch system” refers to a comprehensive system comprising rockets and associated ground infrastructure used to launch spacecraft and payloads into space.
- “launch vehicle” refers to a rocket designed to transport payloads from terrestrial bodies (e.g., Earth, Moon, or Mars) to space or to a designated orbital trajectory.
- “LEO satellite constellation” refers to a network of numerous satellites operating in Low-Earth Orbit, typically deployed to provide services such as broadband connectivity, including Starlink.
- “Low-Earth Orbit” or “LEO” refers to an orbit relatively close to Earth’s surface, typically used by satellites for applications like broadband internet due to its lower latency compared to higher orbits.
- “low-latency network” refers to a network with latency below 70 milliseconds.
- “lunar mass driver” refers to a launch system that we intend to build on the Moon’s surface that will be designed to use electromagnetic acceleration to propel payloads into space without the use of rockets.
- “Macrohard” refers to a platform we are currently developing that is designed to emulate digital workflows, augment human operation of computers, and create a fully AI-operated software company.
- “mass to orbit” refers to the total kilograms of payload deployed to orbit in a given period, and is a key indicator of our capacity and scalability that supports Space revenue and drives expansion across our Connectivity and AI segments.
- “MAU” (or monthly active users) refers to the total number of users who have interacted with Grok or X through web browsers or mobile applications at least once during the 30-day period ending on the date of measurement (“active users”). In presenting combined MAUs across the two platforms, we seek to identify and account for users who access both Grok and X based on sign-in traffic so that such users are not double-counted

[Table of Contents](#)

when measuring MAU. Furthermore, only users who have registered for an X or Grok account are included. While we believe our methodologies provide a reasonable approximation of MAU based on the number of unique users, they may not fully capture all instances of duplication, and our reported MAU should be viewed as an estimate of unique users across our Grok and X platforms for the applicable period. We track the subset of

users who used Grok’s AI features and those who have not based on the source of their server requests.

- “Mbps” refers to megabits per second.
- “Megapack” refers to a containerized, utility-scale lithium-ion battery energy storage system produced by Tesla and designed to stabilize power grids, store renewable energy, and replace fossil fuel peaker plants.
- “megawatt” refers to one million watts.
- “Merlin” refers to the Merlin family of engines, which include vacuum and sea level variants and are fully developed and produced by the Company.
- “microgravity” refers to very weak gravity, such as that experienced in orbiting spacecraft, which allows for unique manufacturing processes like creating ultra-pure materials.
- “Mid-Earth Orbit” or “MEO” refers to an orbital region between approximately 2,000 km and 35,786 km above Earth’s surface.
- “mission success rate” refers to the proportion of Falcon 9 and Falcon Heavy missions that achieve their primary objectives. This term does not include Starship flight tests.
- “mobile network operators” or “MNOs” refers to the local entities of the companies that provide mobile phone services to customers, with whom SpaceX partners to offer satellite-to-mobile connectivity. The term may also include mobile virtual network operators, where applicable.
- “Mobile Satellite Service” refers to providing wireless voice, messaging, and data connectivity to, from, or between mobile devices by using orbiting satellites rather than terrestrial cell towers.
- “Moore’s Law” refers to an observation, not a physical law, that the number of transistors on a microchip doubles roughly every two years, leading to exponentially faster, smaller, and cheaper electronics.
- “orbital AI compute” refers to artificial intelligence computing infrastructure contemplated to be deployed in space, consisting of satellite constellations that act as orbital data centers, harnessing solar energy for power and leveraging the space environment for cooling. We expect to begin deploying our orbital AI compute satellites as early as 2028.
- “payload” refers to the portion of a vehicle’s total mass that consists of the cargo, passengers, satellites, or other mission-specific items being transported and that reaches the target orbit or destination. Payload is distinct from total mass (also referred to as gross mass or initial mass) which is the entire weight of the vehicle, including the payload, fuel / propellant, structure, engines, and any other items, at the start of a journey.
- “payload capacity to orbit” refers to a theoretical payload capacity that a particular launch vehicle is capable of delivering to a specified orbit (e.g., LEO or GEO) or celestial body (e.g., Mars) under specific conditions, which orbit is derived from advanced computer simulations and performance modelling that apply to particular mission scenarios and trajectory assumptions. Actual payload capacity for a given mission may be different and will vary depending on numerous mission parameters and operational factors, including mission-specific trajectory requirements, atmospheric conditions, vehicle and payload configuration, risk profile, and applicable regulatory or range-safety limitations.
- “Power Usage Effectiveness” refers to the global standard metric for data center efficiency, calculated as the ratio of total facility power to IT equipment power.

[Table of Contents](#)

- “propellant” refers to the chemical substance or combination of substances consumed by a rocket engine to produce thrust by generating high-velocity exhaust gases.
- “propulsive landing” refers to the process of landing a rocket or spacecraft using its engines to control descent and achieve a soft, vertical touchdown.
- “radiative cooling” refers to a cooling method that dissipates heat by radiating it into space, often passively, and is expected to be used in orbital AI compute infrastructure.
- “Raptor engines” refers to high-performance family of engines developed and produced by the Company, such as those powering the Super Heavy booster and Starship upper stage, designed for efficiency and reusability.
- “reflight” refers to the reuse of a flight-proven rocket booster or upper stage that has successfully completed a prior space mission, and has been recovered, refurbished, and certified for subsequent launches.
- “return payload mass” refers to the theoretical payload mass that a particular spacecraft is capable of bringing back to Earth from a specified orbit under specific conditions, which is derived from advanced computer simulations and performance modelling that apply to particular mission scenarios and trajectory assumptions. Actual payload that can be returned for a given mission may be different and will vary depending on numerous mission parameters and operational factors, including mission-specific trajectory requirements, atmospheric conditions, vehicle and payload configuration, risk profile, and applicable regulatory or range-safety limitations.
- “rideshare” refers to a type of space mission where multiple satellites or payloads from different customers are launched together on a single rocket, sharing the cost.
- “satellite-to-mobile” refers to a service that provides global cellular connectivity directly to everyday smartphones via satellites, supplementing terrestrial networks and eliminating mobile dead zones.
- “Service Line” refers to an individual instance of Starlink broadband internet service provisioned under a subscription plan, generally associated with a specific Starlink User Terminal or group of terminals, and billed according to Starlink’s service plans and terms of service. The number of Service Lines is distinct from the number of unique devices, account holders, end users, or physical persons.

- “space economy” refers to economic activities related to the development, production, and operation of goods and services that utilize or support space-based infrastructure and capabilities, including launch services, satellite systems, and space-enabled technologies.
- “Space segment” refers to our Space segment, which includes our customer launch operations and offerings such as Falcon, Dragon, and Starship.
- “SpaceX Bridge Loan” refers to the Bridge Loan Credit Agreement, dated as of March 2, 2026, by and among the Company, as borrower, the guarantors from time to time party thereto, the lenders from time to time party thereto and Goldman Sachs Bank USA, as administrative agent and a lender.
- “SpaceX Credit Facility” refers to our Credit Agreement, dated as of February 7, 2025, by and among the Company, as borrower, the guarantors from time to time party thereto, the lenders from time to time party thereto and Bank of America, N.A., as administrative agent, as amended by the First Amendment to Credit Agreement and Waiver, dated as of March 2, 2026, by and among the Company, the lenders party thereto, and the other L/C Issuers party thereto. In May 2026, the SpaceX Credit Facility was amended to increase the borrowing capacity and extends the maturity date.
- “spectrum” refers to the range of electromagnetic frequencies used for wireless communication, with licensed spectrum granting use for specific services.
- “Starlink” refers to our global Low-Earth Orbit satellite constellation and broadband network designed to deliver high-speed, low-latency internet connectivity worldwide.

[Table of Contents](#)

- “Starlink Consumer Broadband” refers to a category of Starlink active users encompassing both individual residential users (households and personal use) and small-to-medium-sized businesses.
- “Starlink Fixed Site” refers to a category of Starlink active users encompassing exclusively enterprise businesses.
- “Starlink Kit” refers to a set of products needed to connect to the Starlink network, typically including a Starlink User Terminal and accessories.
- “Starlink Mobile” refers to a service that provides cellular connectivity directly to everyday smartphones via satellites, supplementing terrestrial networks and substantially reducing mobile dead zones.
- “Starlink Subscriber” refers to a unique Service Line that is directly assigned to a Starlink.com account registered to a person or entity that does not have a direct, negotiated agreement with the Starlink sales team.
- “Starlink User Terminal” refers to a device developed by the Company that connects to the Starlink satellite constellation to deliver high-speed, low-latency internet.
- “Starshield” refers to a secure satellite network designed specifically for government customers and national security applications.
- “Starship” refers to a fully reusable, super heavy-lift launch vehicle. Starship can be used to describe the stacked vehicle (booster and upper stage) or upper stage only. We expect Starship to commence payload delivery to orbit in the second half of 2026.
- “Sun-synchronous orbit” refers to a type of polar orbit around a planet in which a satellite passes over any given point of the planet’s surface at the same local mean solar time, allowing for consistent solar energy capture.
- “Super Heavy” refers to the reusable first-stage booster for the Starship launch vehicle, powered by 33 Raptor engines.
- “SuperGrok” refers to our subscription-based Grok service that provides users with expanded access to Grok models and related tools.
- “SuperGrok Heavy” refers to our subscription-based Grok service tier that provides users with expanded access to Grok models and related tools, including higher usage limits relative to SuperGrok.
- “SuperGrok Lite” refers to our subscription-based Grok service tier that provides users with basic access to Grok models and related tools.
- “supported accounts” refers to, when used in the context of our X platform and Grok, a human, bot or similar account that logged into the X platform or Grok. The total number of supported accounts may include fake, spam or bot accounts if they are active.
- “Tbps” refers to terabits per second.
- “Terafab” refers to a chip manufacturing initiative with a long-term goal of producing one terawatt of compute hardware each year.
- “terawatt” refers to one trillion watts.
- “terawatt-scale” refers to infrastructure, systems, or facilities that are designed to generate, transmit, or consume approximately one terawatt or more of electrical power capacity.
- “terrestrial AI compute” refers to artificial intelligence computing infrastructure located on Earth, such as data centers and supercomputers, used for training and running AI models.

[Table of Contents](#)

- “throughput” refers to the rate at which data or material can be processed or transferred, often referring to network capacity or production output.
- “tokens” refers to the basic units of text or images processed and generated by an AI model, used to measure AI workload, throughput, and computational output.
- “watt” is the International System of Units (SI) unit for measuring power, representing the rate of which energy is transferred, used or generated.
- “X” refers to our real-time information, entertainment, and free speech platform that serves as a foundational distribution and data engine for the AI ecosystem.
- “xAI” refers to X.AI Holdings LLC or, prior to the xAI Merger, X.AI Holdings Corp., together with its subsidiaries, as applicable.
- “xAI Gov” refers to our offering that provides government customers with access to Grok models and related tools for use in governmental applications, workflows, and services.
- “X Premium+” refers to our highest subscription tier for X.

Our Satellite Names

We use a “V” naming convention for our Starlink satellites (such as V1, V2 Mini, and V3). Although we use a similar “V” naming convention for both our broadband and mobile satellite constellations, these are distinct systems. Our broadband satellites are designed to deliver high-speed internet services to homes, businesses, and vehicles, while our mobile satellites are designed to connect directly to cell phones from space. These constellations have different performance requirements and technical specifications. Please see below the terms used for our satellites throughout this prospectus:

- “V1 Broadband satellites” refers to our first-generation satellites in the Starlink constellation for delivering high-speed, low-latency broadband internet from LEO to Earth.
- “V1 Mobile satellites” refers to our mobile satellites that provide light data, text messaging (SMS), and over-the-top voice services (e.g., WhatsApp and FaceTime) to mobile devices. V1 Mobile satellites are currently in orbit and are launched on our Falcon rockets.
- “V2 Mini satellites” refers to our current broadband satellites that provide high-speed internet to homes, businesses, and vehicles. V2 Mini satellites are currently in orbit and are launched on our Falcon rockets.
- “V2 Mobile satellites” refers to our next-generation mobile satellites, which are designed to provide more comprehensive satellite-to-mobile services, including broadband data and IoT connectivity and which we expect to begin deploying on Starship in 2027.
- “V3 satellites” refers to our next-generation Starlink broadband satellites, which are designed to offer one Tbps of downlink capacity per satellite and which we expect to begin deploying on Starship in the second half of 2026.



[Table of Contents](#)

PROSPECTUS SUMMARY

This summary highlights information contained elsewhere in this prospectus. This summary is not complete and does not contain all of the information you should consider before investing in our Class A common stock. You should read this entire prospectus carefully before making an investment decision. You should carefully consider, among other things, the sections titled “Risk Factors,” “Management’s Discussion and Analysis of Financial Condition and Results of Operations,” and our consolidated financial statements and the related notes included elsewhere in this prospectus. Some of the statements in this summary constitute forward-looking statements. Please carefully consider “Cautionary Statement Regarding Forward-Looking Statements.”

“You want to wake up in the morning and think the future is going to be great—and that’s what being a space-faring civilization is all about. It’s about believing in the future and thinking that the future will be better than the past. And I can’t think of anything more exciting than going out there and being among the stars.”

—Elon Musk

Our Mission

Our mission is to build the systems and technologies necessary to make life multiplanetary, to understand the true nature of the universe, and to extend the light of consciousness to the stars. To do this, we have formed the most ambitious, vertically integrated innovation engine on (and off) Earth with unmatched capabilities to rapidly manufacture and launch space-based communications that connect the world, to harness the Sun to power a truth-seeking artificial intelligence that advances scientific discovery, and ultimately to build a base on the Moon and cities on other planets.

Overview

Founded in 2002, SpaceX is the only company building the integrated hardware and software infrastructure of the future across space, connectivity, and AI. At our core, we are builders. We design, manufacture, launch, and operate products and services built on cutting-edge technologies, including the world’s most advanced rockets and spacecraft. We safely and reliably transport astronauts, satellites, and other payloads on missions that benefit life on Earth. Since 2023, we have launched more than 80% of mass to orbit for the world each year with an over 99% mission success rate with Falcon rockets. We also operate a high-speed, low-latency global broadband data and communications network powered by approximately 9,600 Starlink broadband and mobile satellites in Low-Earth Orbit, delivering connectivity to millions of consumer, enterprise, and government customers across 164 countries, territories, and other markets, as of March 31, 2026. Using our dedicated satellite-to-mobile constellation, we offer connectivity services, supplementing terrestrial networks and substantially reducing mobile “dead zones” across approximately 30 countries.

With the potential to improve both space exploration and life on Earth, AI accelerates SpaceX’s mission to make life multiplanetary, to understand the true nature of the universe, and to extend the light of consciousness to the stars. xAI, which was founded in 2023 and acquired by SpaceX in early 2026, is now an integral pillar of our vertically integrated company. We are rapidly constructing AI compute infrastructure—starting on Earth with the goal of extending to space—at industry-leading pace and cost efficiency. Our infrastructure supports training and inference for Grok, which has emerged as one of the world’s most advanced frontier models. Grok is designed as a truth-seeking AI model, built on our founder Elon Musk’s mission to enable humanity to understand the universe. We believe that accomplishing this mission requires a truth-seeking approach to AI. We define truth seeking as the active, relentless pursuit of what is objectively true about reality, and grounded in evidence, logic, empirical data, and first principles thinking. Our goal is to understand and explain what the universe appears to be doing, as

accurately as current knowledge allows. Within two years of its initial model release, Grok achieved frontier-level performance in scientific reasoning, as measured by its GPQA Diamond score, an industry benchmark that evaluates AI models on a standardized set of questions written and validated by experts, on a faster timeline than reported by other leading model providers. Grok also benefits from integration with X, our real-time information, entertainment, and free speech platform, which serves as a foundational distribution and data engine for our AI ecosystem and further enhances Grok’s truth-seeking objective.

[Table of Contents](#)

We believe that space represents the largest economic frontier in human history. Connectivity infrastructure in space is designed to help everyone on Earth have access to education, healthcare, entertainment, and communications, and to enable people to overcome many traditional limits, such as physical and political borders. We believe AI infrastructure in space can utilize the virtually limitless power of the Sun and thereby enable the use of AI as a transformative force for understanding the universe and improving the daily lives of all humans. We believe the convergence of these areas will enable an unprecedented expansion in the global economy, leading to an age of abundance. Our innovations and technological advancements are redefining industries on Earth, while we aim to create new ones on the Moon, Mars, and beyond. We are truly building the infrastructure of the future.

- **Space.** SpaceX is the only company that has cracked the code on accessing space at scale, revolutionizing an industry characterized by decades of stagnation, risk aversion, and economically perverse cost structures. SpaceX upended this paradigm through the application of first-principles thinking, which rejects industry assumptions and builds solutions based on the fundamental laws of physics. Our intense, mission-driven, engineering-first culture and focus on extreme vertical integration have propelled us to achieve what many deemed impossible. We pioneered high-cadence, reliable, and affordable access to space with our Falcon family of rockets. In 2015, we established at least a 10-year lead over the industry by successfully landing our first Falcon 9 booster back from space before anyone else. Space flight that historically cost billions per launch now costs in the tens of millions, fundamentally reducing the cost of space access and providing the opportunity to build new enterprises in space.
- **Connectivity.** Since activating service for customers in 2020, Starlink has rapidly expanded global access to high-speed internet, prioritizing underserved rural and remote communities worldwide. While building terrestrial networks in such communities can be prohibitively expensive, Starlink is capable of delivering broadband connectivity anywhere on Earth with just a Starlink Kit. As of March 31, 2026, we had approximately 9,600 Starlink broadband and mobile satellites in Low-Earth Orbit, operating the world’s most advanced broadband constellation providing internet connectivity to approximately 10.3 million Starlink Subscribers across 164 countries, territories, and other markets. In January 2024, we also began deploying our Starlink Mobile constellation that utilizes separate Starlink satellites with satellite-to-mobile capabilities, substantially reducing mobile “dead zones” around the world. As of March 31, 2026, our dedicated satellite-to-mobile constellation of approximately 650 V1 Mobile satellites provides satellite-to-mobile data, over-the-top voice, and messaging services to approximately 7.4 million monthly unique devices across approximately 30 countries.
- **AI.** We were the first company to deploy a coherent gigawatt-scale AI training cluster. For complex reasoning and agentic workloads, compute is directly correlated with the quality of intelligence and task completion speed. In under two years, we have established a dual advantage in both cost efficiency and deployment speed at scale. By owning the compute infrastructure and vertically integrating across the full AI stack, we can train and iterate our frontier models at lower cost and higher velocity and accelerate development cycles. This eliminates external bottlenecks and drives rapid, continuous improvements in model performance. We believe this combination of our state-of-the-art AI compute infrastructure, our truth-seeking frontier model, and our access to real-time data on X creates a significant strategic advantage. Our integrated AI platforms across Grok and X had approximately 1.3 billion supported accounts active in the last twelve months ended March 31, 2026 and December 31, 2025, including approximately 550 million and 520 million MAUs as of March 31, 2026 and December 31, 2025, respectively. Of our MAUs, we had approximately 117 million and 89 million MAUs that used Grok’s AI features as of March 31, 2026 and December 31, 2025, respectively. Grok’s deep integration with X enables freshness, relevance, and contextual awareness that we believe is a competitive differentiator. This direct, real-time access to the information and human discourse on X enhances Grok’s truth-seeking capabilities by grounding outputs in up-to-date knowledge and diverse viewpoints. As a result, we believe Grok can deliver the most objective and relevant insights and best serve high-frequency, high-value use cases across consumer and enterprise AI applications.

We have created distinct new markets across the space, connectivity, and AI industries by building the integrated hardware and software infrastructure of the future and by combining our broad range of capabilities. For example, SpaceX’s recent acquisition of xAI unites SpaceX’s launch capabilities and global connectivity network with xAI’s AI development capabilities. Specifically, we believe SpaceX’s reusable rockets, scaled satellite manufacturing, and

[Table of Contents](#)

operational expertise can enable the cost-effective and rapid deployment of massive AI compute satellite constellations—with potentially millions of satellites—for orbital data centers. We believe these AI compute satellites in Sun-synchronous orbit will be able to handle energy-intensive AI workloads, such as inference demand, at far greater scale and efficiency than terrestrial alternatives, with Starlink providing low-latency, global connectivity linking these orbital AI systems to people around the world and delivering real-time intelligence. We expect to begin deploying our orbital AI compute satellites as early as 2028.

Our financial results reflect the strength of our operating model and our ability to create and scale multiple new

businesses:

- For the three months ended March 31, 2026, we generated revenue on a consolidated basis of \$4,694 million, loss from operations of \$(1,943) million and Adjusted EBITDA of \$1,127 million. In 2025, we generated revenue on a consolidated basis of \$18,674 million, loss from operations of \$(2,589) million and Adjusted EBITDA of \$6,584 million. Our Space and Connectivity segments contributed the substantial majority of our consolidated revenue in the three months ended March 31, 2026 and the year ended December 31, 2025, demonstrating the benefits of their scale and operating leverage in our vertically integrated business model;
- For the three months ended March 31, 2026, our Space segment generated revenue of \$619 million, loss from operations of \$(662) million, and Segment Adjusted EBITDA of \$(351) million. In 2025, our Space segment generated revenue of \$4,086 million, loss from operations of \$(657) million, and Segment Adjusted EBITDA of \$653 million. Additionally, our Space segment funded \$930 million and \$3,004 million in research and development expense during the three months ended March 31, 2026 and the year ended December 31, 2025, respectively, for our next-generation Starship launch vehicle program. Starship is designed to enable a step-function change in our launch capability across reusability, payload capacity, and launch cadence, and is the key enabler of our long-term growth strategy by unlocking entirely new categories of missions;
- For the three months ended March 31, 2026, our Connectivity segment generated revenue of \$3,257 million, income from operations of \$1,188 million, and Segment Adjusted EBITDA of \$2,087 million. Our Connectivity segment, primarily driven by Starlink, generated revenue of \$11,387 million, income from operations of \$4,423 million, and Segment Adjusted EBITDA of \$7,168 million in 2025, representing year-over-year growth of 49.8%, 120.4%, and 86.2%, respectively, benefiting from subscriber growth, increasing enterprise adoption, and continued improvement in network efficiency;
- In our newly acquired AI segment, we plan to prioritize growth and investment to capture significant opportunities in AI applications and compute infrastructure. For the three months ended March 31, 2026, our AI segment generated revenue of \$818 million, loss from operations of \$(2,469) million, and Segment Adjusted EBITDA of \$(609) million. In 2025, our AI segment generated revenue of \$3,201 million, loss from operations of \$(6,355) million, and Segment Adjusted EBITDA of \$(1,237) million, reflecting its earlier stage of development and continued investments to support long-term growth opportunities in AI; and
- For the three months ended March 31, 2026, capital expenditures for our Space segment was \$1,052 million, for our Connectivity segment was \$1,332 million and for our AI segment was \$7,723 million. In 2025, capital expenditures for our Space segment was \$3,832 million, for our Connectivity segment was \$4,178 million and for our AI segment was \$12,727 million.

Segment Adjusted EBITDA is a non-GAAP measure. Please refer to the section titled “Management’s Discussion and Analysis of Financial Condition and Results of Operations—Non-GAAP Financial Measures” for additional information on our non-GAAP financial measures, including reconciliations of Segment Adjusted EBITDA to segment income (loss) from operations, the most directly comparable GAAP measure.

Why This Matters Now

For the entirety of its existence, human civilization has lived on a single celestial body: Earth. The current paradigm, in which human civilization is confined to one planet, exposes humanity to existential threats that are unpredictable and uncontrollable on a planetary scale. By moving beyond the only home we have ever known, we ensure species-level redundancy and that the light of consciousness will not be tied to a single planet subject to the inevitable

[Table of Contents](#)

hazards of a harsh and vast universe. We do not want humans to have the same fate as dinosaurs. We want to give them a reason to look ahead with excitement, with the prospect that we are entering an age of abundance with an endlessly prosperous and exciting future.

For decades, a reality where humanity travels between the planets and the stars has felt tantalizingly close but still locked in the pages and screens of science fiction. We are capable of better understanding the universe, exploring the universe, and ultimately making life multiplanetary across the universe. We are becoming a civilization with the ability to reach beyond Earth’s cradle and begin to inhabit other worlds. While we remain dedicated to this fundamental mission, our progress in accessing space continues to yield opportunities that enrich life on Earth. For example, by dramatically reducing the cost of access to space, we have been able to expand our mission to address some of the Earth’s most pressing challenges, including bridging the digital divide by aiming to connect over three billion unconnected people to the internet and humanity’s collective knowledge.

The rapid emergence of the AI era intensifies the urgency of our mission, as AI has the potential to accelerate not only space exploration, but also transformative societal advancements on Earth. However, AI’s ability to revolutionize human potential is directly dependent on meeting exponentially increasing resource demands. On Earth, the massive expansion of data center capacity to support growing compute demand is significantly outpacing electricity generation, which was effectively flat in the United States for approximately 15 years, growing at a compound annual growth rate of 0.1% from 2008 to 2023. Despite the recent increase in electricity demand from AI data centers, electricity generation in the United States has grown at an annual rate of less than 3% between 2023 and 2025, while electricity generation in China has grown at approximately twice that rate in the same time period. This supply and demand imbalance is already imposing unsustainable strains on terrestrial power grids, supply chains, and the environment. The Sun contains approximately 99.8% of the solar system’s energy and, as a result, we believe it is the only truly scalable solution to terrestrial energy constraints in the age of AI. Harnessing this energy in space is considerably more efficient than on land. Space-based solar arrays can generate more than five times the energy per unit area of terrestrial solar due to continuous illumination, lack of atmospheric interference, and optimal orientation. SpaceX is well-positioned to capture this space-based solar energy through our ability to rapidly access Sun-synchronous orbit through our satellite manufacturing scale and launch capability. As a result, we are expanding our footprint and harnessing the vast resources of space that are essential to sustaining technological development. Our goal is to ensure that AI becomes a force for human flourishing and a benefit to civilization, rather than a catalyst for terrestrial resource depletion and instability.

We believe that our current space efforts will catalyze transformative breakthroughs that could reshape terrestrial industries and lead to the emergence of new trillion-dollar markets on the Moon, Mars, and beyond. In particular, we believe our goal of establishing a lunar presence will enable terawatt-scale annual AI compute growth, support deeper space exploration and industrialization, and serve as a stepping stone to establishing a civilization on Mars.

We believe that the next paradigm shift for humanity is the creation of a resilient, perpetually expanding spacefaring civilization that drives continuous innovation across new frontiers, ultimately propelling us to Kardashev Type II status—we believe we are capable of unlocking an era of unprecedented economic expansion, while also contributing to the safeguards of humanity’s future against existential risk.

Who We Are

SpaceX combines the most transformative and critical technologies in human history, including reusable rockets, a fully global internet service, satellite-to-mobile communications, a real-time information, entertainment and free speech platform, and a truth-seeking AI system designed to accelerate scientific discovery and augment human capabilities.

Our Unparalleled Launch Capabilities

Since our founding in 2002, SpaceX has cracked the code on accessing space at scale, transforming an industry characterized by decades of stagnation, risk aversion, and economically perverse cost structures. We design, manufacture, launch, and refurbish reusable launch vehicles that provide cost-efficient, reliable, and high-cadence access to space for our own purposes as well as for third-party commercial and government customers. Our

[Table of Contents](#)

extensive vertical integration and end-to-end control over the entire value chain, from design to launch to operations, allows us to achieve unprecedented speed and cost efficiency.

As of March 31, 2026, SpaceX had launched a total mass to orbit of approximately 7,400 metric tons with an over 99% mission success rate across our Falcon rockets. We have completed approximately 650 orbital space launches, and over 540 of those launches were completed by a flight-proven Falcon rocket. With the first successful launch of Falcon 1 in 2008, we became the first private company to successfully launch a liquid-fueled rocket to Earth’s orbit. In December 2015, we achieved what many deemed impossible: landing a rocket launched to space back on Earth. By 2017, we were routinely recovering and reusing the Falcon 9 first-stage booster post-launch, delivering another step-function drop in space access costs via groundbreaking reusability. As of March 31, 2026, our Falcon 9 rockets have demonstrated the ability to refly a first-stage 34 times. With the future deployment of Starship, which is designed to be the world’s first fully and rapidly reusable spacecraft, we aim to reduce the cost to reach orbit by 99% or more relative to the historical average launch cost, establishing the most affordable and scalable path to creating new opportunities in space, such as orbital AI compute and Mars exploration.

Our principal launch vehicles and spacecraft include:

- **Falcon 9.** As the world’s first orbital-class rapidly reusable rocket, Falcon 9 was first launched in 2010 and has a payload capacity to LEO of approximately 23 metric tons when fully expendable. Falcon 9 has completed approximately 620 orbital space launches as of March 31, 2026, and an over 99% mission success rate. According to NASA, the first version of Falcon 9 in 2010 reduced launch cost to approximately \$2,700 per kilogram, approximately 85% less than the historical average launch cost of \$18,500 per kilogram.
- **Falcon Heavy.** Falcon Heavy first launched in 2018 when it put a Tesla all-electric sports car (“Tesla Roadster”) and its mannequin passenger, known as Starman, into orbit around the Sun. With a payload capacity to LEO of approximately 64 metric tons, Falcon Heavy is a partially reusable super heavy-lift launch vehicle designed to deliver large payloads to orbit. Falcon Heavy is one of the most powerful operational rockets in the world measured by liftoff thrust, with 11 launches as of March 31, 2026 and a 100% mission success rate.
- **Dragon.** Launched by Falcon 9 in 2012, our Dragon spacecraft became the first commercial spacecraft to deliver cargo to and from the International Space Station, an orbiting laboratory that serves as a research facility and destination for human spaceflight, and, eight years later, the first privately built vehicle to fly humans to the orbiting laboratory. Since 2020, our Dragon spacecraft has safely flown 78 crewmembers from 20 countries.
- **Starship.** First launched in 2023, Starship is designed to be a fully reusable, super heavy-lift launch vehicle. Starship V3 is designed to deliver 100 metric tons to Earth’s orbit in a fully reusable configuration while enabling rapid turnaround times akin to commercial aviation. Future generations of Starship are being designed to double this payload capacity. To date, we have executed 12 Starship flight tests, with our 12th flight test in May 2026 debuting the next generation Starship vehicle and Super Heavy booster, powered by the next evolution of our Raptor engine and launching from a newly designed pad at Starbase. We expect Starship to commence payload delivery to orbit in the second half of 2026. We have achieved innovative milestones such as catching a booster using “chopstick” arms on the same tower it launched from. We expect this capability will facilitate rapid refurbishment and reuse, allowing for multiple launches per day at reduced costs.

Upon achieving rocket reusability, we recognized the immense potential of our launch business to enable new revenue streams. This led to the development of Starlink, our global satellite internet constellation, consisting of thousands of LEO satellites designed to provide high-speed, low-latency broadband connectivity to underserved areas worldwide. Although the concept of using satellites for global internet connectivity dates back decades, technical challenges and the prohibitive cost of accessing space and deploying the satellites required for capacity and global coverage historically rendered attempts to provide such connectivity economically unviable. Within three years of our first satellite launch in 2019, we solved the technical and production challenges of the satellites and had deployed the largest LEO constellation in existence. Today, Starlink is the sole low-latency network available globally. By combining increasing launch cadence, expanding cargo capacity, and declining unit costs—driven by rapid reusability—we have generated a compounding competitive advantage. This not only fortifies our core business, but also provides vast new market opportunities uniquely enabled by space.

Our Leading Capabilities Across Space, Connectivity, and AI

Space. While our launch capabilities support our other businesses, such as Starlink Consumer Broadband and Starlink Mobile, we also sell launches to third-party customers. We offer launch services to commercial, civil, international and government customers through our reusable Falcon 9 and Falcon Heavy rockets for satellite, cargo, and crew missions. We are the primary launch provider for the U.S. government. In 2025, we launched 11 of 12 National Security Space Launch (“NSSL”) medium and heavy lift missions and all five U.S. crew and cargo missions to the International Space Station for NASA.

Connectivity. Our Connectivity business includes Starlink Consumer Broadband, Enterprise Solutions, Government Solutions, and Starlink Mobile.

- **Starlink Consumer Broadband.** We operate the world’s largest and most advanced space-based internet broadband service. We provide fiber-like download speeds—at a median of 225 Mbps during peak hours for residential users as of March 31, 2026—and the technological capability to provide service everywhere on Earth, including the poles. This service quality is enabled by our vast network of approximately 9,600 Starlink broadband and mobile satellites in Low-Earth Orbit, which accounted for approximately 75% of all active maneuverable satellites in orbit as of March 31, 2026. We expect to commence deploying our next-generation V3 satellites, designed to offer one Tbps of downlink capacity per satellite, using Starship in the second half of 2026. We expect that a single Starship launch will be capable of deploying up to 60 V3 satellites to LEO, representing a potential twenty-fold increase in Starlink downlink capacity deployed relative to a Falcon 9 launch.
- **Enterprise Solutions.** SpaceX is a critical partner to a wide array of enterprises. We offer Starlink’s high-speed, low-latency, reliable internet services to enterprise customers across industries including construction, agriculture, retail, telecom, hospitality, aviation, maritime, and land mobility. Starlink’s unique capabilities are well-suited for deployments across field offices, remote worksites, research stations, drilling rigs, rural hospitals, aircraft, cruise ships, trains, and hotels. We also serve a broad fixed-site customer base across industries such as retail and financial services that require high availability for critical operations as well as reliable connectivity in remote or hard-to-serve locations.
- **Government Solutions.** For our government customers, we provide high-speed, resilient connectivity for public services, social impact, humanitarian efforts, and disaster response in even the most remote and challenging environments. Separately with Starshield, we have leveraged our commercial LEO satellite constellation engineering learnings and operational experiences to develop a secure, dedicated satellite network designed specifically for United States Government customers and national security applications.
- **Starlink Mobile.** We provide satellite-to-mobile connectivity, supplementing terrestrial networks and substantially reducing mobile “dead zones” across approximately 30 countries. Through our partnerships with approximately 30 MNOs on six continents, we enable consumers, businesses, and public-sector customers to use their existing phones in more places, support critical connectivity during disasters and power outages, and open new applications for low-bandwidth mobile and IoT devices.

AI. We operate a highly vertically integrated AI platform.

- **AI Compute Infrastructure.** xAI has established a leading position in building and scaling terrestrial AI compute infrastructure, becoming the first company to deploy a coherent gigawatt-scale AI training cluster. We own and operate what we believe to be the largest AI training data center clusters on Earth, including COLOSSUS and COLOSSUS II. The addition of Terafab, a chip manufacturing initiative with Tesla and Intel, aims to further extend our vertical integration to chip design and manufacturing to alleviate potential future chip shortages at SpaceX, optimize compute performance, and potentially reduce overall compute costs. In connection with such collaboration, we have agreed with Tesla on a general framework for the future development of Terafab. Any specific projects undertaken pursuant to this framework will be subject to separate negotiations and agreements (including any development timelines, milestones and capital expenditures) and have not yet been determined. We believe that the key constraints in the continued growth of AI are physical—

[Table of Contents](#)

chip manufacturing, data center infrastructure, and power generation; the future of AI will be determined by the control of the physical stack.

- **Truth-Seeking Frontier Model.** Since launching Grok-1 in November 2023, we have released four major versions and notable variations thereof, achieving one of the fastest iteration cycles in the industry. Within two years of its initial model release, Grok achieved frontier-level performance in scientific reasoning, as measured by its GPQA Diamond score, an industry benchmark that evaluates AI models on a standardized set of questions written and validated by experts, on a faster timeline than reported by other leading model providers. Building on this trajectory, we expect to continue scaling Grok through subsequent generations. Ongoing training of next-generation models is expected to scale toward multiple trillions of parameters, which could represent a step change in reasoning in depth and overall intelligence. In this context, the number of parameters refers to the scale of the model, where parameters are the internal numerical values, such as “weights,” that are adjusted during training to enable the model to recognize patterns and relationships in data. A larger number of parameters generally allows the model to capture more complex relationships, store greater amounts of knowledge, and achieve higher levels of reasoning capability. This accelerated rate of innovation stems from our highly vertically integrated stack: full ownership of training infrastructure; access to the world’s most powerful compute clusters; and relentless focus on truth seeking and real-world utility. A key competitive differentiator is Grok’s deep integration with X, enabling proprietary access to a real-time information stream of approximately 350 million daily posts, which enhances freshness, relevance, and contextual awareness for Grok. This direct, real-time access to the information and human discourse on X enhances Grok’s truth-seeking capabilities by grounding outputs in up-to-date knowledge and diverse viewpoints.
- **Consumer and Enterprise Applications.** We leverage our leading frontier models and compute infrastructure to deliver consumer and enterprise applications. Together with Tesla, we are also developing Macrohard, an

agentic AI platform designed to be capable of fully emulating digital workflows and augmenting human operation of computers using sophisticated autonomous agents. We believe Macrohard will have the potential to fundamentally transform how companies are structured and operate, thereby allowing dramatic increases in human productivity.

Our Repeatable Business Model

Our business model is built on a repeatable, engineering-driven framework that combines our unparalleled launch capabilities, extreme vertical integration, rapid iteration, and disciplined capital investment to create durable, large-scale businesses. We execute this framework through the following core principles:

1. Leverage our unparalleled launch capabilities to enable massive scale;
2. Identify and create new trillion-dollar market opportunities;
3. Design a solution with world-class engineering and first-principles thinking;
4. Apply “The Algorithm” (make less dumb, delete, optimize, accelerate, automate);
5. Vertically integrate all the way to the end customer;
6. Continuously drive cost down and throughput up; and
7. Generate significant cash flow and reinvest in the future.

Our Engineering-First Culture

We are able to achieve transformative technological breakthroughs because we accept only the laws of physics as the limiting factors to our work and mission. Our core approach is deeply rooted in first-principles thinking, which rejects any preconceived notions or experience-based norms. We have a track record of achieving what many have deemed impossible. Some of our industry-defining achievements and historic milestones include:

- The first private company to develop and launch a liquid-fuel rocket to reach orbit (2008);
- The first private company to successfully dock a private spacecraft with the International Space Station (2012);
- The first to successfully propulsively land (2015) and reflly orbital-class rocket boosters (2017);
- The first to begin deploying a large-scale LEO broadband satellite constellation (2019);
- The first private company to transport astronauts to orbit, returning America’s ability to fly astronauts to and from the International Space Station (2020);

[Table of Contents](#)

- The first to manufacture consumer-grade phased-array user terminals at scale (2022);
- The first to deploy a large-scale LEO satellite-to-mobile constellation (2025);
- The first to build a gigawatt-scale AI training cluster and largest coherent supercomputer (2026);
- The first gigawatt-scale Megapack battery installation (2026); and
- The only company capable of building orbital AI compute at scale.

Our AI Compute Infrastructure Advantage and Growth Strategy

Why Compute Matters. We believe AI leadership will be defined by the ability to rapidly scale compute capacity to support exponential usage growth and frontier intelligence. The training and inference demanded by advanced AI models require substantial computational resources. Reasoning models introduced in 2024 demonstrated that allocating more computational resources and giving models more time to process during inference directly leads to higher-quality intelligence. In addition, compute infrastructure with end-to-end, cluster-level coherence through tight integration across software and hardware systems enables more efficient, stable, and higher-fidelity training and inference at scale—ultimately enhancing model intelligence and performance. Within inference, we expect computationally-intensive reasoning, agentic, and multi-modal workloads will continue to grow as a portion of overall usage. We therefore believe operators with superior model-to-compute integration—the ability to efficiently support and allocate compute across both training and inference workloads—are best positioned to win the AI race.

Self-Reinforcing Network Effects Among Lower Cost Per Token, Model Quality, and User Adoption. AI systems are ultimately constrained or differentiated by the cost, speed, and scale at which they can generate and process tokens. A “token” represents the fundamental unit of data consumed and produced by modern AI models. This is because lower cost per token enables more frequent model training, larger and more sophisticated models, longer chains of processing for reasoning and agentic workloads, and significantly higher inference volumes at economically viable prices. This dynamic directly impacts model quality, responsiveness, and accessibility, while also determining the ability to serve the rising global demand across consumer, enterprise, and mission-critical AI applications. This creates a self-reinforcing advantage in which lower token costs drive greater model quality and user adoption, reinforcing AI leadership.

Cost of Compute is the Main Driver of Cost Per Token. The total cost per token is determined by the efficiency, availability, and unit economics of the underlying compute and the cost of building and operating compute infrastructure. Improvement in the cost of building and operating this compute infrastructure—whether through lower data center construction cost, lower power infrastructure cost, shorter time to grid interconnection, or higher cluster-level throughput—translates directly into lower cost per token. Accordingly, for a given level of intelligence, we expect the long-term economics of AI companies to be driven by the ability to consistently deliver bleeding-edge compute at the lowest possible cost per token. Put simply, we view cost per token as a function of three primary inputs—the underlying AI model, the compute hardware, and energy, and we expect to have a competitive advantage in the latter two cost components. We believe we have a pathway over time that will significantly reduce compute hardware costs through continued vertical integration and development of proprietary chips, building on our experience designing custom silicon for our Starlink satellites. We also expect that the marginal cost of energy for our AI compute satellites will be minimal because our satellites are powered by solar arrays in space. By driving the energy component to minimal levels and pursuing improvements in compute hardware cost, we believe we can achieve a meaningfully lower overall cost per token in the future.

We Have a Dual Speed and Cost Advantage in Terrestrial AI Compute. We own and operate what we believe to be the largest AI training data center clusters on Earth. Our AI compute facilities, COLOSSUS and COLOSSUS II, collectively provide approximately 1.0 gigawatt of compute power, with additional power capacity available for data center operations. Our first-principles thinking enables us to build coherent compute at scale and at rapid speed with lower costs than most other companies in the industry. In order to bring compute clusters online as fast as possible, we employ a vertically integrated, nimble approach to construction. We brought the first cluster of COLOSSUS online in 122 days, repurposing the shell of an existing factory, and the first cluster of COLOSSUS II online even faster in 91 days. As an illustrative comparison, an industry benchmark to bring online a 100 megawatt greenfield

[Table of Contents](#)

We Believe Orbital AI Can Accelerate Time to Power and Reduce Token Costs. The Sun contains approximately 99.8% of the solar system's energy and offers what we believe is the only truly scalable solution to the challenge of accelerating demand for compute relative to terrestrial energy constraints. The logical path forward is to move power-intensive AI workloads into orbit, where solar energy is near-constant and uninterrupted. With such accessibility to energy, we believe that our launch business will enable us to consistently activate the highest performing hardware before our competitors without such access, shrinking the timeline to useful tokens on bleeding-edge hardware and sustaining our token cost advantage. We believe SpaceX is uniquely positioned to deploy and operate data centers in orbit that can eventually achieve a lower cost than terrestrial data centers over time due to our extreme vertically integrated approach across launch, satellite manufacturing at scale, network connectivity, and terrestrial data center expertise.

We Believe We Are Well-Positioned to Deliver Orbital AI Compute. We believe orbital AI compute is an incredibly difficult technical challenge that only we can solve at scale in the near term. We are the only company that has already accomplished the key technical challenges associated with evolving connectivity satellites into AI compute satellites. In our view, we are well-positioned to deliver a full-scale AI compute satellite constellation. Significant work remains, but we are confident in our singular leadership position.

- **We have unmatched satellite launch capabilities to enable deployment at scale.** Deployment of 100 gigawatts per year via satellites carrying over 100 kilowatts of compute power per metric ton will require thousands of launches per year and the transport of approximately one million metric tons to orbit annually. The fully reusable nature of Starship positions us to be capable of launching this level of mass. Starlink Broadband V1 and V2 Mini satellites have already demonstrated launch survivability and high reliability under vibration, shock, g-loads, acoustic stress, and vacuum exposure, achieving 99.9% average uptime.
- **We have already solved many of the significant technical hurdles to evolving connectivity satellites into AI compute satellites.** Through our leading expertise of connectivity satellites—including mass production, deployment, network operations, and inter-satellite lasers and mesh connectivity—we have already solved the hardest part in the development of AI compute satellites. Because AI compute satellites represent an evolution of spacecraft engineering already demonstrated through Starlink, we believe development of AI compute satellites will be easier for us than for anyone else. Our existing Starlink constellation is another crucial enabler of orbital AI compute, as its global network allows data from our AI compute satellites to reach ground stations anywhere on Earth.
- **We will use our proven Starlink in-orbit technology to optimize our orbital AI compute.** In order to operate orbital AI compute satellites, we plan to build on our vast experience of operating approximately 9,600 Starlink broadband and mobile satellites in Low-Earth Orbit. In 2025 alone, Starlink satellites proactively performed over 1,000 automated collision avoidance maneuvers per day guided by this technology to safely and efficiently operate the constellation. This operating model gives us control over workload placement across Earth and space while maintaining resilience through redundancy and fail safe systems. A high degree of controllability will allow the satellite to be optimized for brightness mitigation, disposal, and other modes of operation.
- **We can manufacture our AI compute constellations at scale with rapid upgrade cycles.** We have built one of the largest satellite manufacturing operations in the world. Our vertically integrated approach with limited reliance on third-party suppliers will be key to our mass-scaling efforts and should allow us to deploy the latest AI processors. We believe SpaceX will be the first and only company to manufacture satellites at the scale of automotive manufacturing.
- **We are building chip manufacturing capabilities to scale our access to AI compute hardware.** We announced a collaboration with Tesla in March 2026 to build the Terafab initiative with a long-term goal of producing one terawatt of compute hardware each year. In connection with such collaboration, we have agreed with Tesla on a general framework for the future development of Terafab. Intel joined the project in April 2026 and is expected to contribute its expertise in designing, fabricating, and packaging ultra-high performance chips to help Terafab scale. Any specific projects undertaken pursuant to this framework will be subject to separate negotiations and agreements (including any development timelines, milestones and capital expenditures) and

[Table of Contents](#)

have not yet been determined. With this internal manufacturing capability, we plan to alleviate potential future chip shortages at SpaceX, especially as we develop orbital AI at scale, and design chips that are optimized for the space environment.

- **We can leverage our terrestrial experience to build and operate compute clusters and AI workloads at scale.** We believe our experience operating compute infrastructure on Earth provides the technical and operational foundation to extend these capabilities into orbit. For example, we plan to subject compute hardware to extensive pre-deployment testing on Earth to identify early life failures before launch to reduce in-orbit disruption. For compute hardware that does fail, we plan to leverage existing Starlink fleet management software to reallocate traffic to other satellites and prevent cluster-level downtime.

We Believe Our Infrastructure is a Distinct Advantage in Delivering Superior AI. We expect the combination of

competitive cost per token, our ability to deploy and operate data centers in orbit, and our strength in connectivity to result in more scalable intelligence that is accessible globally at high speeds.

Our Strengths

- Global Leadership in Orbital Launch Services
- Unrivaled Satellite and Connectivity Platform across Design, Manufacturing, Deployment, and Operations
- Truth-Seeking AI Model Enhanced by Real-Time Data
- Extreme Vertical Integration Enabling High Velocity and Superior Cost Efficiency at Scale
- Unique Ability to Scale New Trillion-Dollar Markets Across Space, Connectivity, and AI
- Business Models that Are Incredibly Difficult to Replicate
- Mission-Driven Culture and World-Class Talent

Our Growth Strategies

Space

- Increase launch payload capacity
- Establish the lunar economy, including cargo transport, manufacturing, and energy production on the Moon

Connectivity

- Grow Starlink Consumer Broadband and enterprise and government customers
- Expand our Starlink Mobile offering
- Increase the capacity of our constellations

AI

- Grow consumer AI platform monetization
- Grow X monetization
- Deepen enterprise and government adoption
- Increase the scale of our terrestrial power and AI compute infrastructure
- Deploy orbital AI compute at scale
- Design and manufacture our own chips
- Launch digital human augmentation

Future Markets


- Point-to-point terrestrial travel
- Space tourism
- In-orbit manufacturing
- Passenger and cargo transport to the Moon and Mars
- Energy production on the Moon and Mars
- Manufacturing capabilities on the Moon and Mars
- Asteroid mining

[Table of Contents](#)

Our Market Opportunity

We believe we have identified the largest actionable total addressable market (“TAM”) in human history. We estimate that our quantifiable TAM is \$28.5 trillion, consisting of \$370 billion in Space from space-enabled solutions; \$1.6 trillion in Connectivity across \$870 billion in Starlink Broadband and \$740 billion in Starlink Mobile as well as additional opportunities in enterprise and government; \$26.5 trillion in AI across \$2.4 trillion in AI infrastructure, \$760 billion in consumer subscriptions, \$600 billion in digital advertising, and \$22.7 trillion in enterprise applications. For illustrative purposes of sizing our addressable market opportunity, we exclude China and Russia from our global estimates.

SpaceX's Estimated TAM by Segment



Our Challenges

We face a number of challenges relating to our business and growth strategy and, ultimately, the achievement of our mission to make life multiplanetary, understand the true nature of the universe, and extend the light of consciousness to the stars. The pursuit of our mission drives our decision-making and forms the foundation of our business plan,

which is predicated on building, commercializing, and operating services and products at a scale that has not previously been achieved. This objective requires us to develop and integrate complex and novel technologies, develop new processes and infrastructure, and coordinate across multiple suppliers, contractors, regulators, and stakeholders. Because we are attempting to execute at a scale for which there is no precedent, we face heightened uncertainty with respect to design, engineering, procurement, construction, commissioning, and operational performance. In particular, our ability to execute our growth strategy is highly dependent on the successful development and scaling of Starship and the ability to increase our launch cadence, both of which are subject to challenges and uncertainties inherent in the development and deployment of new and complex technologies. Additionally, many of our initiatives described above under “Our Growth Strategies,” including those to develop orbital AI compute at scale, manufacture AI chips at scale, establish a lunar economy, transport humans and cargo to the Moon and Mars, and develop human augmentation systems, involve significant technical complexity, unproven technologies or technologies that do not exist, and such initiatives may not achieve commercial viability. Many of the innovative products and services described elsewhere in this prospectus may ultimately be unsuccessful and may require great expense, innovations not yet achieved or technologies not yet developed. As a result, the timeline for certain of our initiatives involving unproven or new innovations, including our goal of deploying 100 gigawatts of

[Table of Contents](#)

annual compute power to orbit, the establishment of a lunar economy and interplanetary industrialization, and the launch cadence required to achieve these goals may be difficult or impossible to determine. Our growth strategy may take longer to execute than anticipated, and you may not realize a return on your investment within the timeframe you anticipate, or at all.

In addition, a portion of our anticipated market opportunities is associated with industries described above under “Future Markets.” Certain of these industries, such as space tourism and cargo transport to the Moon, are still emerging. Others, including in-orbit manufacturing, passenger transport to the Moon, passenger and cargo transport to Mars, energy production on the Moon and Mars, manufacturing capabilities on the Moon and Mars, and asteroid mining, do not exist today. While we believe these industries will develop over time, the manner in which they emerge, including the timing of commercialization, the scale and pace of adoption, and the applicable competitive, technical, regulatory, geopolitical, and economic frameworks may differ materially from our current expectations.

Our Space, Connectivity, and AI segments are also subject to the following challenges and uncertainties, among others.

- **Space:** Our growth strategy depends on our ability to increase our launch cadence and payload capacity, which is dependent on the successful development of Starship at scale. Unexpected design modifications, supply chain disruptions, anomalies, environmental issues, and other unforeseen technical challenges could result in delays or failures to deploy Starship on our anticipated schedule, which would delay or impede our ability to achieve our other business objectives, such as the deployment of our next-generation satellites, the expansion of our satellite-to-mobile connectivity services, and deployment of in-orbit AI compute infrastructure.
- **Connectivity:** Our satellite connectivity, including our global satellite-to-mobile connectivity services under Starlink Mobile, depend on access to radio frequency spectrum and authorizations from the Federal Communications Commission (the “FCC”) in the United States and telecommunications regulators in other countries. Acquiring the necessary authorizations can be a complex and time-consuming process. Without these licenses and approvals, we cannot generally offer connectivity services in a given market. Spectrum access itself is limited and highly regulated. Additionally, the growth of our connectivity services depends on our ability to increase market awareness and acceptance of connectivity through Starlink across numerous international markets, each with its unique challenges.
- **AI:** Our AI business is in a relatively early stage, it is being integrated into our organization, its business strategy is still developing, and it will require significant capital expenditures to fund compute, infrastructure and power generation, model training, and product development. Additionally, our AI business is subject to challenges inherent in a nascent, highly competitive, capital intensive and rapidly changing industry. These include the potential for disruptive technological change, evolving industry and regulatory standards, the emergence of new and well-funded competitors, frequent new product and service introductions, and changing customer demands.

Any number of these challenges, and others that may be currently unknown to us, could have a negative impact on our business, financial condition, and results of operations. For a discussion of the challenges, risks, and limitations that could harm our future prospects, please refer to “Cautionary Note Regarding Forward-Looking Statements,” “Risk Factors,” and “Management’s Discussion and Analysis of Financial Condition and Results of Operations” included elsewhere in this prospectus.

Recent Developments

Collaboration with Cursor

In April 2026, we entered into a compute and option agreement with Anysphere, Inc., doing business as Cursor, a San Francisco-based private software company (“Cursor”), which we view as a compelling extension of our strategy to vertically integrate compute infrastructure, models, and applications. Under the compute agreement, we will provide Cursor with certain GPU cluster compute capacity and collaborate to improve existing models, including Grok, and potentially to jointly develop AI models and related model-specific deliverables or products. With the option agreement, we have the right, but not obligation, to acquire Cursor at a predetermined price or pay a fee. We

[Table of Contents](#)

consider software development as a strategically important use case for AI given its combination of high-quality structured data, rapid feedback cycles and frequent, mission-critical usage. AI-assisted coding workflows generate context-rich, verifiable data that can enhance model training and performance, while also driving sustained inference demand. The depth of Cursor’s integration with a high-frequency coding workflow generates valuable developer interaction data, including coding generation prompts, iteration cycles, and software architecture decisions. We expect that access to this data will enhance our model training and inference, including with respect to Grok. Meanwhile, by providing access to our large-scale compute infrastructure, we believe we can help Cursor deliver faster and higher quality user experiences. The collaboration with Cursor may also accelerate our AI strategy by integrating our AI models more directly into developer workflows and expanding the distribution of our AI capabilities through high-engagement software interfaces.

The consideration for the acquisition of Cursor, if any, after the closing of this offering would consist of shares of our Class A common stock based on an implied equity value of Cursor of \$60.0 billion, and the price of our Class A common stock that equals the volume-weighted average closing price thereof over the seven consecutive trading days immediately preceding the closing of the acquisition. If either (i) we decide to terminate the option agreement or (ii) Cursor is eligible to and decides to terminate due to our material breach of the option agreement (subject to notice and cure provisions), Cursor is entitled to a \$1.5 billion termination fee under the option agreement and an \$8.5 billion deferred services fee under the compute agreement. These fees are payable in cash (or Class A common stock, if this offering has not been consummated at the time the fees become payable). For more information about our arrangement with Cursor, including our option to acquire the company, please refer to “Business—Collaboration with Cursor” included elsewhere in this prospectus.

Compute Services Agreements with Third Parties

We believe our compute infrastructure and related strategy provides us with substantial flexibility in how we allocate and monetize capacity. We have the ability to use compute resources to support our proprietary AI applications (such as Grok-5, which is currently being trained at COLOSSUS II), while also providing access to select compute capacity to third-party customers. For example, in May 2026, we entered into Cloud Services Agreements with Anthropic PBC (“Anthropic”), an AI research and development public benefit corporation, with respect to access to compute capacity across COLOSSUS and COLOSSUS II. Compute capacity provided includes approximately 325,000 NVIDIA GPUs, backed by hyperscale-class CPUs, exabyte-scale storage and high-speed networking and interconnects purpose-built for AI workloads. Pursuant to these agreements, the customer has agreed to pay us \$1.25 billion per month through May 2029, with capacity ramping in May and June 2026 at a reduced fee. After the initial three-month period, the agreements may be terminated by either party upon 90 days’ notice. The customer will retain ownership and intellectual property rights in its content, AI models, and related data. This structure allows us to monetize a portion of the compute capacity in our infrastructure, while still permitting reallocation of that capacity for our own internal initiatives if needed in the future. We have sufficient capacity to provide compute for our own AI models, including support of our training and inference demands, and to satisfy the obligations under these agreements. We expect to enter into additional similar services contracts. We believe this opportunity highlights the increasing importance of large-scale, frontier-level AI infrastructure and positions us as a differentiated provider of high-performance compute capacity to both internal and third-party AI workloads. We believe our dual monetization strategy provides multiple pathways to generate returns on invested capital.

Founder, Chief Executive Officer, Chief Technical Officer and Chairman of Our Board

Mr. Musk is our founder, Chief Executive Officer, Chief Technical Officer and the Chairman of our board. Assuming a size as set forth on the cover page of this prospectus and an initial public offering price of \$135.00 per share, Mr. Musk will hold approximately 82.4% of the voting power of our common stock (or 82.3% if the underwriters exercise their option to purchase additional shares of Class A common stock in full) immediately after this offering through his ownership of 849,494,440 shares of our Class A common stock and 5,219,053,075 shares of our Class B common stock, which comprises approximately 91.6% of our Class B common stock. The foregoing amount of Class B common stock includes 1,302,072,285 restricted shares of Class B common stock issued to and held of record by Mr. Musk, which may be voted by Mr. Musk, and the vesting of which is subject to the satisfaction of certain performance and other conditions. Under our charter, the holders of our Class B common stock will have the right to elect a majority of our board (such directors, the “Class B Directors”), for so long as any shares of Class

[Table of Contents](#)

B common stock remain outstanding. As the holder of a majority of our shares of Class B common stock, Mr. Musk will be able to elect, remove or fill any vacancy among the Class B Directors. In addition, for so long as he beneficially owns more than 50% of the voting power of our common stock, Mr. Musk will control the voting power over the selection of our board. As a result, Mr. Musk will have the power to control the outcome of matters requiring shareholder approval, including election of all our directors, and to control our business and affairs.

Our Controlled Company Status

We will be a controlled company as of the completion of this offering under Nasdaq and Nasdaq Texas listing rules. A controlled company is not required to have a majority of its board composed of independent directors or to establish independent compensation and nominating committees. As a controlled company, we will remain subject to rules that require us to have an audit committee composed entirely of independent directors.

Corporate Information

We were founded and incorporated as Space Exploration Technologies Corp., a Delaware corporation, on March 14, 2002 and reincorporated as a Texas corporation on February 14, 2024. Our principal executive offices are located at 1 Rocket Road, Starbase, Texas 78521. Our website address is www.spacex.com. Information contained on our website or linked therein or otherwise connected thereto does not constitute part of nor is it incorporated by reference into this prospectus or the registration statement of which this prospectus forms a part.